

Koshiro Aoki

Master's student in Computer Science, Waseda University

✉ aokikoshiro@akane.waseda.jp

🌐 aokikoshiro.github.io

🐙 AokiKoshiro

🌐 koshiro-aoki

Research Interests

Metacognition | Introspection | (Mechanistic) interpretability | NeuroAI

Education

Waseda University

M.S. in Computer Science
Adviser: Daisuke Kawahara

Apr. 2026 - Present

📍 Tokyo, Japan

Waseda University

B.S. in Computer Science
GPA: 3.8/4. Dean's Award.

Apr. 2022 - Mar. 2026

📍 Tokyo, Japan

Work Experience

Waseda University

Research Assistant

Research on differences between the internal representations of the human brain and language models as part of a JST CREST project.

Apr. 2026 - Present

📍 Tokyo, Japan

The University of Tokyo

Research Contractor

Research on interpretability and defenses of adversarial attacks on robot foundation models.

May 2025 - Present

📍 Tokyo, Japan

Research and Development Center for Large Language Models, National Institute of Informatics (NII)

Technical Assistant

Research on bottom-up interpretability of training dynamics in large language models.

May 2025 - Mar. 2026

📍 Tokyo, Japan

Publications

Conference

In-Context Neurofeedback: Can Large Language Models Control Their Internal Representations through Privileged Access?

2026

Koshiro Aoki, Ryota Takatsuki, Gouki Minegishi, Yusuke Haruki, Daisuke Kawahara

The 2026 Conference on Empirical Methods in Natural Language Processing (EMNLP 2026 Findings)

Trustworthy AI for Good (AI4GOOD) Workshop @ ICML 2026 Spotlight

Bottom-Up Interpretability of Pretraining Dynamics via Loss Curve Decomposition

2026

Koshiro Aoki, Masaru Isonuma, Max Müller-Eberstein, Yusuke Oda, Hirokazu Kiyomaru, Takashi Kodama, Chaoran Liu, Yohei Oseki, Yusuke Miyao, Daisuke Kawahara

Third Conference on Language Modeling (COLM 2026)

Workshop

Crosscoders Identify Shared or Specific Features between the Human Brain and Language Models

2026

Koshiro Aoki, Itsuki Hamada, Naho Orita, Daisuke Kawahara, Hiromu Sakai

Mechanistic Interpretability Workshop @ ICML 2026

Evaluating the Impact of SAE-based Language Steering on LLM Performance

2026

Sebastian Zwirner, Wentao Hu, Koshiro Aoki, Daisuke Kawahara

Testing Simulation Theory in LLMs' Theory of Mind	2025
Koshiro Aoki , Daisuke Kawahara	
<i>IJCNLP-AACL 2025 Student Research Workshop</i>	Oral

Domestic Journal

Mechanistic Interpretability: A New Trend in Interpretability Research	2026
Koshiro Aoki , Ryota Takatsuki, Gouki Minegishi, Hakaze Cho	
<i>Transactions of the Japanese Society for Artificial Intelligence</i>	Best Paper Award

Honors and Awards

Spotlight Paper	2026
<i>Trustworthy AI for Good (AI4GOOD) Workshop @ ICML 2026</i>	
Best Paper Award	2026
<i>The Japanese Society for Artificial Intelligence 40th Anniversary Commemorative Paper Award</i>	
Dean's Award for Excellence	2026
<i>School of Fundamental Science and Engineering, Waseda University</i>	
Young Researcher Encouragement Award	2026
<i>The 32nd Annual Meeting of the Association for Natural Language Processing</i>	
Oral Presentation	2025
<i>IJCNLP-AACL 2025 Student Research Workshop</i>	
Encouragement Award	2025
<i>The 20th Young Researcher Association for NLP Studies</i>	
Sponsored Award	2025
<i>The 31st Annual Meeting of the Association for Natural Language Processing</i>	